



MASTER DE MATHÉMATIQUES
PARCOURS "Modélisation et Analyse Numérique"

PROGRAMMATION II - HAX011X - 2025/2026

HW - sparse_matrix

A naive implementation of matrices (arrays) is suitable for dense matrices, of which most elements are nonzero. For sparse matrices, on the other hand, a much more efficient approach is to store only the nonzero elements and ignore all the zero elements. This approach makes arithmetic operations much more efficient and saves a lot of time and storage.

Sparse matrices are thus a most useful tool in practical implementations. The data structures used in their implementation, however, should not be simple arrays anymore but more elaborated containers, like for instance `linked_lists`.

The items `T` in the `linked_list` should be objects that contain two fields: one to store the actual value of the element and the other to store the index of the column in which it lies. The sparse matrix should thus be implemented as a `list` of such `linked_lists`.

In summary, we need three new objects, which are studied below:

- the `row_element` object, which contains data about a particular element in a given row;
- the `row` object, which is a `linked_list` of `row_element`s objects and models a given row of the matrix;
- the `sparse_matrix` object, which is a `list` of `rows` objects.

Exercice 1. The `row_element` object

- (1) The `row_element` class contains two **protected** data fields. The first field `value_` is of type `T`, to be specified later in compilation time. This field contains the value of the element. The second field `column_` is of type `int` and contains the index of the column in which the element is located in the matrix.
- (2) As usual, although the data fields are protected, they can still be read/set by the public `value()` and `column()` member functions.
- (3) The constructor should have two arguments corresponding to the two protected fields. Default values should be provided: 0 for `value_`, -1 for `column_`. Of course -1 is a meaningless value which should be updated during the construction of the matrix.
- (4) The class should have a copy constructor and a definition of `operator=`.
- (5) Next, some useful arithmetic operations should be defined. For example, the `operator+=()` can take either a `T` or a `row_element` argument. In either case, its `value_` field is incremented by the corresponding value of the argument (note that this `operator+=()` will be used in the `operator+=()` of the `linked_list` merging algorithm, in order to add coefficients). In a similar way, `operator-=`, `operator*=` and `operator/=` should be defined.

(6) Next, binary `operator<()`, `operator>()`, and `operator==()` as nonmember operators should be introduced. These operators take two `row_element` arguments and compare their `column_` fields. For example, `operator<` returns 1 if the column index of the first argument is smaller than that of the second one, and 0 otherwise (this priority order will be used later to preserve increasing column order in matrix rows). Note that this `operator==()` will be used in the `operator+=()` of the `linked_list` merging algorithm, in order to check if coefficients have the same column before adding their value_.

(7) We should also define some useful nonmember binary arithmetic operators that involve a `row_element` object and a scalar. For instance:

```
const row_element<T> operator+(row_element<T> const& e, T const& t)
```

should return a `row_element` with updated `value_` field. here is the required list:

```
const row_element<T> operator+(row_element<T> const& e, T const& t);
const row_element<T> operator+(T const& t, row_element<T> const& e);
const row_element<T> operator-(row_element<T> const& e, T const& t);
const row_element<T> operator-(T const& t, row_element<T> const& e);
const row_element<T> operator*(row_element<T> const& e, T const& t);
const row_element<T> operator*(T const& t, row_element<T> const& e);
const row_element<T> operator/(row_element<T> const& e, T const& t);
```

(8) Finally, a function that prints a `row_element` to the screen is needed.

Test your class in a main program.

Exercice 2. The row object

In this Exercice, we introduce the `row` class that implements a row in a sparse matrix. The `row` object is actually derived from `linked_list<row_element<T>>`.

The `row` class is also a template class. The template `T` that indicates the type of value of the element is to be specified later in compilation time. It is assumed that the elements in the `row` object are ordered in increasing order. The priority order of `row_element` objects used for this purpose is induced from the priority order of their corresponding columns.

(1) the constructor should be defined as follows:

```
row(T const& val=0, int col=-1){
    item=row_element<T>(val,col);
}
```

so that it calls the constructor of the `row_element` class, and insert the corresponding instance of `row_element` as the first element of the list (note that the default constructor of the base class `list<row_element>` is automatically called in priority when this constructor is invoked, creating an empty list, which will then be populated through methods of the derived class),

(2) a definition of the operator `operator()()`, declared as follows

```
row_element<T> const& operator()() const;
```

should return the first item of the list only,

(3) a method declared as follows

```
T const& value() const;
```

should return the first-item value,

(4) a method `int column() const` should return the first-item column index.

Since the `row` class is derived from the `linked_list` class, it can use its public and protected members. In particular, it can use the functions that insert or drop items. Still, the `row` class should contains its own local versions of the `insert_next_item()`, `insert_first_item()`, and `append()` functions. Indeed, these versions

should be different from the original versions coming from the class `linked_list`, since they take two arguments to specify the `value_` and `column_` of the inserted element.

- (5) For instance, here is the definition of the `insert_next_item` method

```
void insert_next_item(T const& val, int col){
    row_element<T> e(val, col);
    linked_list<row_element<T>>::insert_next_item(e);
} // insert a row_element as second item
```

Note that this new version calls the method from the base class. Based on this example, define modified versions of the `insert_first_item` and `append` methods,

- (6) Define a function `const T row_sum() const` using recursion in order to compute the sum of the elements in the row.

We have seen that the recursive structure of the base `linked_list` class is particularly useful. The definitions of many member functions use recursive calls applied to the `p_next_` field that points to the rest of the row. However, the `p_next_` field inherited from the base `linked_list` class is of type pointer-to-`linked_list` rather than pointer-to-`row`. Therefore, it must be converted explicitly to pointer-to-`row` before the recursive call can take place. This is done by adding the prefix (`row*`). Note that usually, this conversion is considered risky because in theory `p_next_` can point to a `linked_list` object or any object derived from it, with a completely different interpretation of the recursively called function. Fortunately, here `p_next_` always point to a `row` object, so the conversion is safe.

- (7) Recursion may also be used to define again `operator[]`. This function should take an integer argument `i` and returns a copy of the value of the element in column `i`, if it exists. If, on the other hand, there is no such element in the current row, then it returns 0.

This should be carried out using the recursive pattern of the `row` object. First, the `column_` field in the first element is examined. If it is equal to `i`, then the required element has been found, and its value is returned as output. If, on the other hand, it is greater than `i`, then there is no hope of finding the required element because the elements are ordered in increasing column order, so 0 is returned. Finally, if it is smaller than `i`, then the `operator[]` is applied recursively to the rest of the row.

As before, the `p_next_` field must be converted explicitly from pointer-to-`linked_list` to pointer-to-`row` before recursion can be applied to it. Again, this is done by adding the prefix (`row*`).

Note also that the value returned by the `operator[]` function is of type `constant-T` rather than reference-to-`constant-T`. This is because, as discussed above, the function may also return the zero value. This local constant cannot be referred to and must be stored in the temporary unnamed variable returned by the function.

- (8) Recursion should also be used in the following arithmetic operators that involve a `row` and a scalar:

```
const row& operator*=(T const& t);
const row& operator/=(T const& t);
```

- (9) Fortunately, the `operator+=` that adds a `row` to the current `row` (while preserving increasing column order) is already available from the base `linked_list` class, so there is no need to write it again. However, we have to adapt it to possibly *add* the values of the coefficients (remember the remarks about it during PW2, when writing the merging algorithm):

```
// hey : overloaded comparison operator==() for item_ (actually: row) may be used
if (p_scan2 && p_scan1->item() == p_scan2->item())
{
    // hey : overloaded comparison operator+=(()) for item_ (actually: row) may be used
    p_scan1->item() += p_scan2->item();
    p_scan2 = p_scan2->p_next();
}
```

- (10) A binary operator that computes the inner product of a `row` and a `dynamic_vector` is also required. This operation will be used later to compute the product of a sparse matrix and a vector. It may be defined as follows:

```
const T operator*(dynamic_vector<T> const& v) const;
```

- (11) Also required is a function that re-numbers the columns with new numbers contained in a vector of integers named `renumber`. To increase efficiency, this vector is passed to the function by reference. As before, recursion is applied to the `p_next_` field after its type is converted from pointer-to-linked_list to pointer-to-row. The corresponding declaration is:

```
void renumber_columns(dynamic_vector<int> const& renumber);
```

- (12) We also need to define some non-member binary arithmetic operators that involve a row and a scalar:

```
const row<T> operator*(row<T> const& r, T const& t);
const row<T> operator*(T const& t, row<T> const& r);
const row<T> operator/(row<T> const& r, T const& t);
```

- (13) Finally, we offer the following `drop_items` function, which takes as argument a vector of integers named `mask`. Happy Christmas ! The zeroes in this vector indicate that the row elements in the corresponding columns should be dropped. This is done by looking ahead to the next element and dropping it if appropriate. For this purpose, the `row` object that contains the rest of the elements in the row is first accessed as `*p_next_`. Then, the first item in this `row` object (which is actually the second element in the current row) is accessed as `(*next)()`, using the `operator()` in the base `linked_list` class. Now, the column of this element can be read, and if the corresponding component in `mask` vanishes, then this element is dropped by the `drop_next_item` function of the base `linked_list` class. This leads to the following code:

```
template<class T>
void row<T>::drop_items(dynamic_vector<int> const& mask){
    if (p_next()){
        if (!mask[(*p_next()).column()])
            drop_next_item();
        drop_items(mask);
    }
    else
        (*row<T>*p_next()).drop_items(mask);
    if (!mask[column()])
        drop_first_item();
}
// "masking" the row by a vector of integers
```

Exercice 3. The sparse_matrix object

In this Exercice, we introduce the `sparse_matrix` class. By implementing a sparse matrix as a list of `linked_lists` or, more specifically, as a list of `row` objects, only the nonzero matrix elements are stored and participate in calculations, whereas the zero matrix elements are ignored. Although `linked_lists` have their own drawbacks in terms of efficiency (because they use indirect indexing that may slow down the performance due to more expensive data access), this drawback is far exceeded by the advantage of avoiding trivial calculations. Furthermore, in some cases it is possible to map the `linked_list` to a more continuous data structure and make the required computations in it. The hierarchy of objects used to implement the sparse matrix is as follows:

- the `sparse_matrix` object, in the highest level, is implemented as a list of `rows` objects,
- the `row` object is by itself implemented as a `linked_list` of `row_elements` objects,
- the `row_element` object, at the lowest level, contains a template parameter `T` to store the value of the element.

Hence, the `sparse_matrix` class is a `template` class derived from a list of `row<T>` objects. Therefore, it enjoys access to the `public` and `protected` members of the `list` class. The additional member functions are defined below in the `sparse_matrix` class. In these methods, you may often loop over all the items

in the underlying list of rows. In this loop, member functions of the row class are used to access the row_elements.

Your new sparse_matrix class should include:

- (1) a first (default) constructor `sparse_matrix(int Nrows=0)` which proceeds to the memory allocation for the list of `row<T>` according to the value of `Nrows` (precisely: allocate an array `row<T>*[Nrows]` and set all the pointers values to zero),
- (2) a second constructor declared as `sparse_matrix(int Nrows, T const& a);` which deals with the memory allocation for the list `<row<T>>` according to the value of `Nrows`, and set the i^{th} row to be `item_[i] = new row<T>(a,i);`. (this leads to a scalar matrix stored as a `sparse_matrix`: only the diagonal terms are stored),
- (3) a third constructor, allowing to directly read a `sparse_matrix` from a provided file, in a formatted way like the *Harwell-Boeing* file format, should be provided for practical use and left as an Exercice (not required for the HW to be validated),
- (4) a definition of `const T operator()(int i, int j)` that returns the $(i, j)^{\text{th}}$ matrix element (if this element is of course non-zero),
- (5) functions that returns the number of rows in the matrix, the number of columns, and the order (actually `max(row_number(), column_number())`) of a non-square matrix, declared as follows:

```
int row_number() const;
int column_number() const;
int order() const;
```

- (6) some member functions dedicated to arithmetic operations:

```
const sparse_matrix& operator+=(sparse_matrix<T> const&);
const sparse_matrix& operator-=(sparse_matrix<T> const&);
const sparse_matrix& operator*=(T const&);
```

The first two ones add (resp. subtract) a `sparse_matrix` to (resp. from) the current `sparse_matrix`. The third one multiplies the current `sparse_matrix` by a scalar,

- (7) some `friend` free functions declared as follows:

```
template<class S>
friend const sparse_matrix<S> operator+(sparse_matrix<S> const& M1, sparse_matrix<S> const& M2);
template<class S>
friend const sparse_matrix<S> operator-(sparse_matrix<S> const& M1, sparse_matrix<S> const& M2);
template<class S>
friend const sparse_matrix<S> operator*(sparse_matrix<S> const& M, const S& t);
template<class S>
friend const sparse_matrix<S> operator*(const S& t, sparse_matrix<S> const& M);
```

- (8) a `sparse_matrix × dynamic_vector` product function. The returned value should be a `dynamic_vector` of suitable size. The function should be declared `friend` as follows:

```
template<class S>
friend const dynamic_vector<S> operator*(sparse_matrix<S> const& M, dynamic_vector<S> const& v)
```

Note that the structure of the `sparse_matrix` as a list of rows should be used. The product method of the `list` class (declared as `operator*(dynamic_vector<T> const& v) const`) can be useful,

- (9) a `sparse_matrix × sparse_matrix` product should also be defined, introduced as follows:

```
friend const sparse_matrix<T> operator*(sparse_matrix<T> const&, sparse_matrix<T> const&);
```

In order to take benefit from the data structure of the `sparse_matrix`, a product algorithm that relies on rows operations should be implemented. Specifically, we observe that if A is a matrix with N rows and B is a matrix with N columns, with B_i the i^{th} row of B , then the i^{th} row of BA

can be written as a linear combination of rows in A with coefficients in b_i as follows:

$$(BA)_i = B_i A = \sum_j B_{ij} A_i.$$

- (10) a function that computes the transposition of a `sparse_matrix`, declared as follows:

```
friend const sparse_matrix<T> transpose<T>(sparse_matrix<T> const&);
```

- (11) a function that extract the diagonal of a `sparse_matrix` and return it as another `sparse_matrix` of same size, declared as follows:

```
template<class T>
const sparse_matrix<T> diagonal(sparse_matrix<T> const& M);
```

- (12) a print free function declared as follows:

```
template<class S>
friend void print(sparse_matrix<S> const& M);
```

which prints on screen the non-zero coefficients, together with their row-column indices,

- (13) a printf free function declared as follows:

```
template<class S>
friend void printf(sparse_matrix<S> const& M, ostream& os);
```

which prints the non-zero coefficients row, column and value in an "out-file-stream" for plotting purpose (see the next Exercice below).

Test your class in a main program: show me with many examples that your hierarchy of classes is correctly operating.